

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv('boston.csv')
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7

Next steps: [View recommended plots](#)

```
df.isnull().sum()
```

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV     0
dtype: int64
```


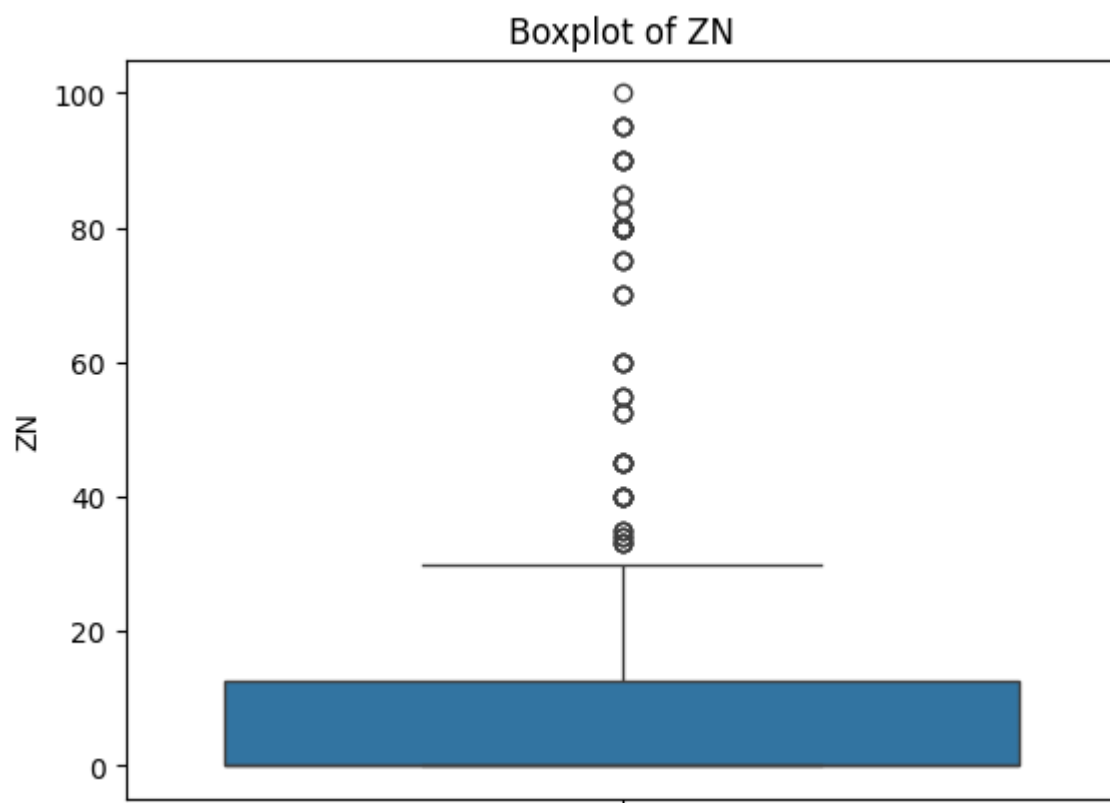
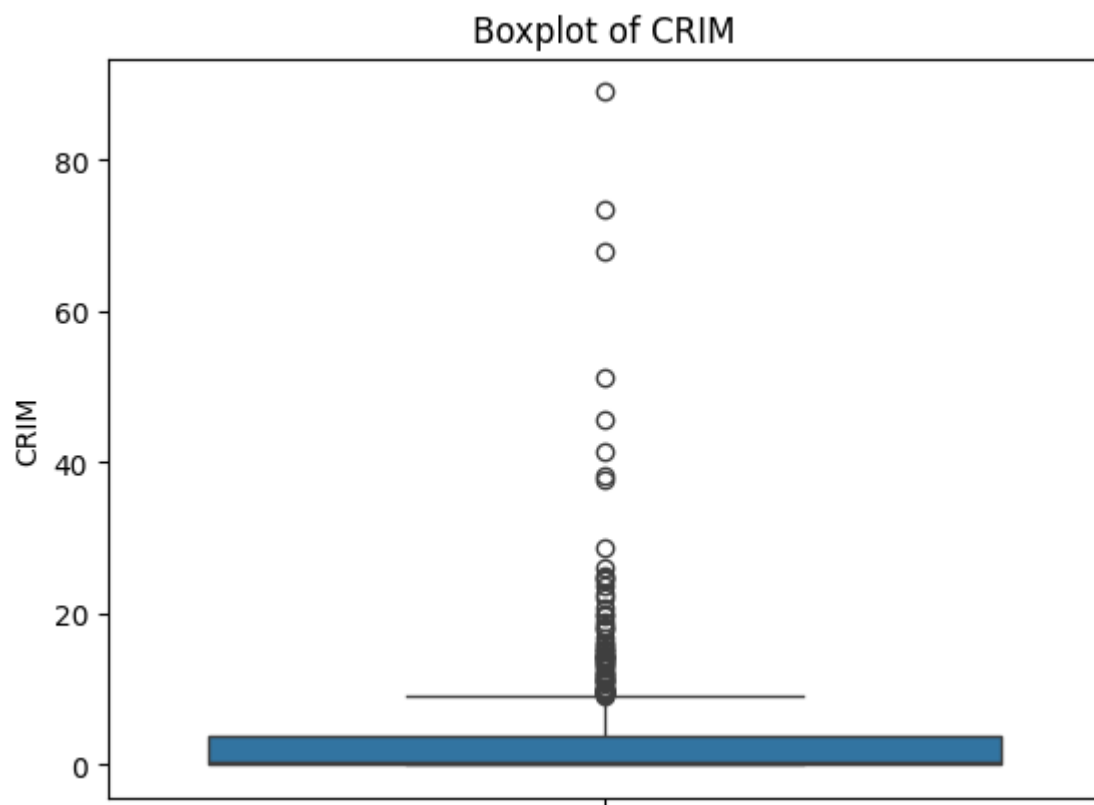
```
numerical_features = []
categorical_features = []
```


```
for f in df.columns:
    if df[f].dtype != 'O':
        numerical_features.append(f)
    else:
        categorical_features.append(f)
```

```
print(numerical_features)
print(categorical_features)
```


```
['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'I
[]
```

```
for f in df.columns:  
    sns.boxplot(df[f])  
    plt.title(f'Boxplot of {f}')  
    plt.show()
```







Boxplot of CHAS




Boxplot of NOX







Boxplot of RM




Boxplot of AGE




Boxplot of DIS







Boxplot of RAD




Boxplot of TAX







Boxplot of PTRATIO




Boxplot of B



Boxplot of LSTAT







Boxplot of MEDV



```
numericalCols = df.select_dtypes(include='number').columns
for f in numericalCols:
    q3 = df[f].quantile(0.75)
    q1 = df[f].quantile(0.25)
    iqr = q3 - q1
    upper_bound = q3 + (1.5 * iqr)
    lower_bound = q1 - (1.5 * iqr)
    df.loc[df[f] >= upper_bound, f] = upper_bound
    df.loc[df[f] <= lower_bound, f] = lower_bound
```

```
for f in df.columns:  
    sns.boxplot(df[f])  
    plt.title(f'Boxplot of {f}')  
    plt.show()
```







Boxplot of CHAS




Boxplot of NOX







Boxplot of RM




Boxplot of AGE




Boxplot of DIS







Boxplot of RAD




Boxplot of TAX







Boxplot of PTRATIO




Boxplot of B




Boxplot of LSTAT






Boxplot of MEDV



```
df.hist(figsize=(24,24), bins=100)
```

```
array([[<Axes: title={'center': 'CRIM'}>, <Axes: title={'center': 'ZN'}>,
       <Axes: title={'center': 'INDUS'}>,
       <Axes: title={'center': 'CHAS'}>],
      [<Axes: title={'center': 'NOX'}>, <Axes: title={'center': 'RM'}>,
       <Axes: title={'center': 'AGE'}>, <Axes: title={'center': 'DIS'}>],
      [<Axes: title={'center': 'RAD'}>, <Axes: title={'center': 'TAX'}>,
       <Axes: title={'center': 'PTRATIO'}>,
       <Axes: title={'center': 'B'}>],
      [<Axes: title={'center': 'LSTAT'}>,
       <Axes: title={'center': 'MEDV'}>, <Axes: >, <Axes: >]],
      ..
      ..
      ..)
```



```
df.skew()
```

```

    CRIM      1.282313
    ZN        1.261340
    INDUS     0.295022
    CHAS      0.000000
    NOX       0.729308
    RM        0.296640
    AGE      -0.598963
    DIS       0.908467
    RAD       1.004815
    TAX       0.669956
    PTRATIO  -0.762495
    B        -1.164208
    LSTAT     0.808671
    MEDV      0.353614
    dtype: float64

```

```
# Model building
```

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error

```

```

X = df.drop('MEDV', axis=1)
y = df['MEDV']

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
regr = LinearRegression()
regr.fit(X_train, y_train)

```

```

▼ LinearRegression
LinearRegression()

```

```

# predicting the model output
y_pred = regr.predict(X_test)

```

```

mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

```

```
Mean Squared Error: 13.264821903384027
```

```

print("Tuned Weights:")
for i in range(len(X.columns)):
    print(f"{X.columns[i]}: {regr.coef_[i]}")
print(f"Intercept: {regr.intercept_}")

```

```

Tuned Weights:
CRIM: -0.47702088562571165
ZN: 0.04146279314123744
INDUS: -0.021440789997329568
CHAS: 8.340550472496489e-14
NOX: -14.041056617657127
RM: 2.7458448668745974
AGE: -0.0027275908048000872

```

```
ASE: 0.002722998870000072  
DIS: -1.0342233144791397  
RAD: 0.28211799925832115  
TAX: -0.008092199582586938  
PTRATIO: -0.8728299586706918  
B: 0.006769390573664431  
LSTAT: -0.4916902018584183  
Intercept: 37.97381451281785
```